

# Decision Tree vs Random Forest in R

```
# Disable warning messages globally  
options(warn = - 1)
```

Library for data wrangling

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
##   filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

Library for timing

```
library(tictoc)
```

Library for decision tree

```
library(rpart)  
library(rpart.plot)
```

Library for random forest

```
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:dplyr':  
##  
##   combine
```

## Library for plotting

```
library(ggplot2)
```

```
##
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:randomForest':
##
##     margin
```

## Importing the data

```
data = read.csv('C:\\Users\\Dell\\OneDrive\\College_2nd\\R Lab\\adult.csv')
```

## Looking at the data

```
head(data)
```

```
##   age workclass fnlwgt   education education.num marital.status
## 1  90      ?  77053    HS-grad           9      Widowed
## 2  82 Private 132870    HS-grad           9      Widowed
## 3  66      ? 186061 Some-college        10      Widowed
## 4  54 Private 140359    7th-8th           4      Divorced
## 5  41 Private 264663 Some-college        10      Separated
## 6  34 Private 216864    HS-grad           9      Divorced
##      occupation relationship race    sex capital.gain capital.loss
## 1      ? Not-in-family White Female           0         4356
## 2 Exec-managerial Not-in-family White Female           0         4356
## 3      ? Unmarried Black Female           0         4356
## 4 Machine-op-inspct Unmarried White Female           0         3900
## 5 Prof-specialty Own-child White Female           0         3900
## 6 Other-service Unmarried White Female           0         3770
##   hours.per.week native.country income
## 1         40 United-States <=50K
## 2         18 United-States <=50K
## 3         40 United-States <=50K
## 4         40 United-States <=50K
## 5         40 United-States <=50K
## 6         45 United-States <=50K
```

## Statistical summary

```
summary(data)
```

```
##      age      workclass      fnlwgt      education
## Min.   :17.00  Length:32561  Min.    : 12285  Length:32561
## 1st Qu.:28.00  Class :character  1st Qu.: 117827  Class :character
## Median :37.00  Mode  :character  Median : 178356  Mode  :character
## Mean   :38.58                      Mean    : 189778
## 3rd Qu.:48.00                      3rd Qu.: 237051
## Max.    :90.00                      Max.    :1484705
## education.num marital.status  occupation  relationship
## Min.      : 1.00  Length:32561  Length:32561  Length:32561
## 1st Qu.: 9.00  Class :character  Class :character  Class :character
## Median :10.00  Mode  :character  Mode  :character  Mode  :character
## Mean     :10.08
## 3rd Qu.:12.00
## Max.     :16.00
##      race      sex      capital.gain  capital.loss
## Length:32561  Length:32561  Min.      : 0  Min.      : 0.0
## Class :character  Class :character  1st Qu.: 0  1st Qu.: 0.0
## Mode  :character  Mode  :character  Median : 0  Median : 0.0
##                      Mean   : 1078  Mean   : 87.3
##                      3rd Qu.: 0  3rd Qu.: 0.0
##                      Max.    :99999  Max.    :4356.0
## hours.per.week native.country  income
## Min.      : 1.00  Length:32561  Length:32561
## 1st Qu.:40.00  Class :character  Class :character
## Median :40.00  Mode  :character  Mode  :character
## Mean     :40.44
## 3rd Qu.:45.00
## Max.     :99.00
```

### Structure of the data

```
str(data)
```

```
## 'data.frame': 32561 obs. of 15 variables:
## $ age : int 90 82 66 54 41 34 38 74 68 41 ...
## $ workclass : chr "?" "Private" "?" "Private" ...
## $ fnlwgt : int 77053 132870 186061 140359 264663 216864 150601 88638 422013 70037
## ...
## $ education : chr "HS-grad" "HS-grad" "Some-college" "7th-8th" ...
## $ education.num : int 9 9 10 4 10 9 6 16 9 10 ...
## $ marital.status: chr "Widowed" "Widowed" "Widowed" "Divorced" ...
## $ occupation : chr "?" "Exec-managerial" "?" "Machine-op-inspct" ...
## $ relationship : chr "Not-in-family" "Not-in-family" "Unmarried" "Unmarried" ...
## $ race : chr "White" "White" "Black" "White" ...
## $ sex : chr "Female" "Female" "Female" "Female" ...
## $ capital.gain : int 0 0 0 0 0 0 0 0 0 ...
## $ capital.loss : int 4356 4356 4356 3900 3900 3770 3770 3683 3683 3004 ...
## $ hours.per.week: int 40 18 40 40 40 45 40 20 40 60 ...
## $ native.country: chr "United-States" "United-States" "United-States" "United-States"
## ...
## $ income : chr "<=50K" "<=50K" "<=50K" "<=50K" ...
```

### Checking for Nulls

```
any(is.na(data))
```

```
## [1] FALSE
```

Converting to factors:

```
data$workclass <- factor(data$workclass, exclude = c("", NA))

data$education <- factor(data$education, exclude = c("", NA))

data$education.num <- factor(data$education.num, exclude = c("", NA))

data$marital.status <- factor(data$marital.status, exclude = c("", NA))

data$occupation <- factor(data$occupation, exclude = c("", NA))

data$relationship <- factor(data$relationship, exclude = c("", NA))

data$race <- factor(data$race, exclude = c("", NA))

data$sex <- factor(data$sex, exclude = c("", NA))

data$income <- factor(data$income, exclude = c("", NA))

data$native.country <- factor(data$native.country, exclude = c("", NA))
```

Looking at structure again

```
str(data)
```

```
## 'data.frame': 32561 obs. of 15 variables:
## $ age : int 90 82 66 54 41 34 38 74 68 41 ...
## $ workclass : Factor w/ 9 levels "?","Federal-gov",...: 1 5 1 5 5 5 8 2 5 ...
## $ fnlwgt : int 77053 132870 186061 140359 264663 216864 150601 88638 422013 70037
## ...
## $ education : Factor w/ 16 levels "10th","11th",...: 12 12 16 6 16 12 1 11 12 16 ...
## $ education.num : Factor w/ 16 levels "1","2","3","4",...: 9 9 10 4 10 9 6 16 9 10 ...
## $ marital.status: Factor w/ 7 levels "Divorced","Married-AF-spouse",...: 7 7 7 1 6 1 6 5 1
## 5 ...
## $ occupation : Factor w/ 15 levels "?","Adm-clerical",...: 1 5 1 8 11 9 2 11 11 4 ...
## $ relationship : Factor w/ 6 levels "Husband","Not-in-family",...: 2 2 5 5 4 5 5 3 2 5
## ...
## $ race : Factor w/ 5 levels "Amer-Indian-Eskimo",...: 5 5 3 5 5 5 5 5 5 5 ...
## $ sex : Factor w/ 2 levels "Female","Male": 1 1 1 1 1 1 2 1 1 2 ...
## $ capital.gain : int 0 0 0 0 0 0 0 0 0 0 ...
## $ capital.loss : int 4356 4356 4356 3900 3900 3770 3770 3683 3683 3004 ...
## $ hours.per.week: int 40 18 40 40 40 45 40 20 40 60 ...
## $ native.country: Factor w/ 42 levels "?","Cambodia",...: 40 40 40 40 40 40 40 40 1 ...
## $ income : Factor w/ 2 levels "<=50K", ">50K": 1 1 1 1 1 1 1 2 1 2 ...
```

Dropping unwanted columns


```
# Drop variables
df <- data %>% select(-c(fnlwgt, capital.gain))
```

```
glims(df)
```

```
## Rows: 32,561
## Columns: 13
## $ age          <int> 90, 82, 66, 54, 41, 34, 38, 74, 68, 41, 45, 38, 52, 32, ~
## $ workclass    <fct> ?, Private, ?, Private, Private, Private, Private, Stat~
## $ education    <fct> HS-grad, HS-grad, Some-college, 7th-8th, Some-college, ~
## $ education.num <fct> 9, 9, 10, 4, 10, 9, 6, 16, 9, 10, 16, 15, 13, 14, 16, 1~
## $ marital.status <fct> Widowed, Widowed, Widowed, Divorced, Separated, Divorce~
## $ occupation   <fct> ?, Exec-managerial, ?, Machine-op-inspct, Prof-specialt~
## $ relationship <fct> Not-in-family, Not-in-family, Unmarried, Unmarried, Own~
## $ race         <fct> White, White, Black, White, White, White, White, White, ~
## $ sex          <fct> Female, Female, Female, Female, Female, Female, Male, F~
## $ capital.loss  <int> 4356, 4356, 4356, 3900, 3900, 3770, 3770, 3683, 3683, 3~
## $ hours.per.week <int> 40, 18, 40, 40, 40, 45, 40, 20, 40, 60, 35, 45, 20, 55, ~
## $ native.country <fct> United-States, United-States, United-States, United-Sta~
## $ income       <fct> <=50K, <=50K, <=50K, <=50K, <=50K, <=50K, <=50K, >50K, ~
```

## Plotting Education


```
ggplot(df, aes(education)) +
  geom_bar(aes(fill=factor(education)), alpha=0.5)
```



We can see that most of them are high school graduates


## Plotting Marital status

```
ggplot(df,aes(marital.status))+  
  geom_bar(aes(fill=factor(marital.status)),alpha=0.5)
```




## Plotting Occupation

```
ggplot(df,aes(occupation))+  
  geom_bar(aes(fill=factor(occupation)),alpha=0.5)
```



## Plotting Race

```
ggplot(df,aes(race))+
  geom_bar(aes(fill=factor(race)),alpha=0.5)
```




The data is clearly unbalanced and biased

## Plotting sex

```
ggplot(df,aes(sex))+  
  geom_bar(aes(fill=factor(sex)),alpha=0.5)
```






Again we can see that the classes are imbalanced

Plotting income

```
ggplot(df,aes(income))+  
  geom_bar(aes(fill=factor(income)),alpha=0.5)
```



The target class is highly imbalanced

## Data pre process for Modelling

Create test and train

```
create_train_test <- function(data, size = 0.8, train = TRUE)
{
  #' create_train_test(df, size = 0.8, train = TRUE)
  #' arguments:
  #' @param df: Dataset used to train the model.
  #' @param size: Size of the split. By default, 0.8. Numerical value
  #' @param train: If set to `TRUE`, the function creates the train set, otherwise the test s
  et. Default value sets to `TRUE`. Boolean value. You need to add a Boolean parameter beca
  use R does not allow to return two data frames simultaneously.

  #' @return test/train data

  n_row = nrow(data)
  total_row = size * n_row
  train_sample <- 1: total_row
  if (train == TRUE) {
    return (data[train_sample, ])
  } else {
    return (data[-train_sample, ])
  }
}
```

## ##Getting data

```
data_train <- create_train_test(df, 0.8, train = TRUE)
data_test <- create_train_test(df, 0.8, train = FALSE)
dim(data_train)
```

```
## [1] 26048    13
```

## Seeing propotion of data

```
prop.table(table(data_train$income))
```

```
##
##      <=50K      >50K
## 0.7470439 0.2529561
```

```
prop.table(table(data_test$income))
```

```
##
##      <=50K      >50K
## 0.8077691 0.1922309
```

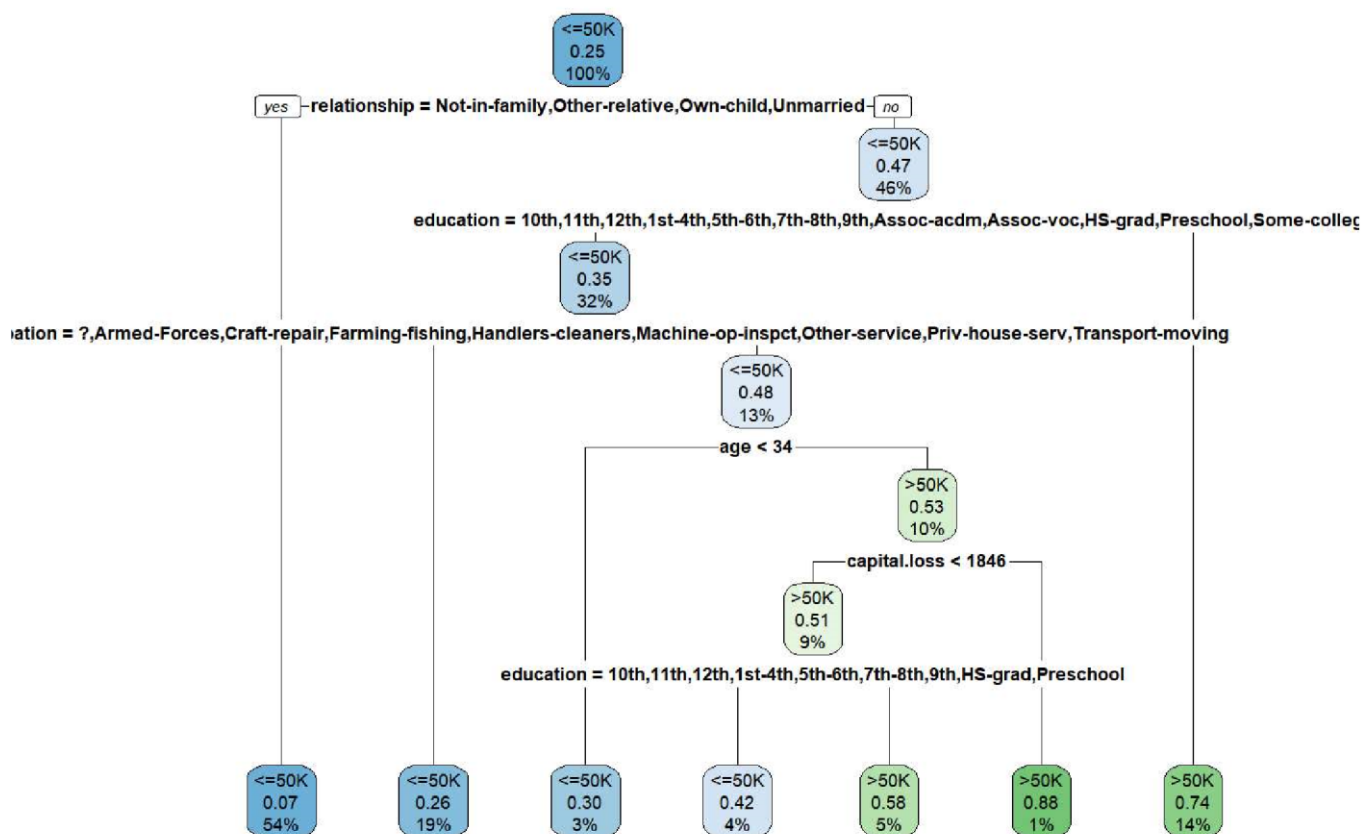
## Decision Tree

```
tic("Running the decision tree: ")
fit_dt <- rpart(income~., data = data_train, method = 'class')
```

```
## Running the decision tree: : 0.8 sec elapsed
```

Plotting the tree,

```
rpart.plot(fit_dt, extra = 106)
```



```
predict_unseen <- predict(fit_dt, data_test, type = 'class')
```

```
table_mat <- table(data_test$income, predict_unseen)
table_mat
```

```
##      predict_unseen
##      <=50K >50K
## <=50K  4859  402
## >50K   573  679
```

## Evaluations on Decision tree

```
accuracy_Test <- sum(diag(table_mat)) / sum(table_mat)
```

```
n = sum(table_mat) # number of instances
nc = nrow(table_mat) # number of classes
diag = diag(table_mat) # number of correctly classified instances per class
rowsums = apply(table_mat, 1, sum) # number of instances per class
colsums = apply(table_mat, 2, sum) # number of predictions per class
p = rowsums / n # distribution of instances over the actual classes
q = colsums / n # distribution of instances over the predicted classes
```

```
precision = diag / colsums
recall = diag / rowsums
f1 = 2 * precision * recall / (precision + recall)
```

## Printing Metrics for Decision tree

```
print(table_mat)
```

```
##          predict_unseen
##          <=50K >50K
## <=50K  4859  402
## >50K   573  679
```

```
print(paste(accuracy_Test, "is the accuracy"))
```

```
## [1] "0.850299401197605 is the accuracy"
```

```
print(paste(precision, "is the precision"))
```

```
## [1] "0.894513991163476 is the precision" "0.628122109158187 is the precision"
```

```
print(paste(recall, "is the recall"))
```

```
## [1] "0.923588671355256 is the recall" "0.542332268370607 is the recall"
```

```
print(paste(f1, "is the f1"))
```

```
## [1] "0.908818853455532 is the f1" "0.582083154736391 is the f1"
```

## Random Forest

### Making the model

```
tic("Running random Forest: ")
model_rf <- randomForest(income ~ ., data = data_train, importance = TRUE)
toc()
```

e

```
## Running random Forest: : 67.87 sec elapsed
```

```
predict_unseen <- predict(model_rf, data_test, type = 'class')
```

```
predict_unseen = as.data.frame(predict_unseen)
```

```
table_mat <- table(data_test$income, predict_unseen$predict_unseen)
```

## Evaluations on Random Forest

```
accuracy_Test <- sum(diag(table_mat)) / sum(table_mat)
```

```
n = sum(table_mat) # number of instances
nc = nrow(table_mat) # number of classes
diag = diag(table_mat) # number of correctly classified instances per class
rowsums = apply(table_mat, 1, sum) # number of instances per class
colsums = apply(table_mat, 2, sum) # number of predictions per class
p = rowsums / n # distribution of instances over the actual classes
q = colsums / n # distribution of instances over the predicted classes
```

```
precision = diag / colsums
recall = diag / rowsums
f1 = 2 * precision * recall / (precision + recall)
```

```
print(table_mat)
```

```
##
##          <=50K >50K
## <=50K    4810  451
## >50K      505  747
```

```
print(paste(accuracy_Test, "is the accuracy"))
```

```
## [1] "0.853216643635805 is the accuracy"
```

```
print(paste(precision, "is the precision"))
```

```
## [1] "0.904985888993415 is the precision" "0.623539232053422 is the precision"
```

```
print(paste(recall, "is the recall"))
```

```
## [1] "0.914274852689603 is the recall" "0.596645367412141 is the recall"
```

```
print(paste(f1, "is the f1"))
```

e

```
## [1] "0.909606656580938 is the f1" "0.609795918367347 is the f1"
```

# INFERENCE:

First, let us see the score of random forest and decision tree

The different evaluations that were done

## Decision Tree :

1. accuracy -> 85.03 | This means that the decision tree has correctly predicted the class on the test data 94.09% of the time
2. precision -> 89.45 | The model gave correct predictions for a class 1, the model predicted correctly
3. precision -> 62.81 | The model gave correct predictions for a class 2, the model predicted correctly
4. recall -> 0.9 | This is the fraction of instances of a class 1 that were correctly predicted, that is 0.9
5. recall -> 0.5 | This is the fraction of instances of a class 2 that were correctly predicted, that is 0.9
6. f1 -> 90.88 | This is the harmonic mean of precision and recall, for class 1
7. f1 -> 58.21 | This is the harmonic mean of precision and recall, for class 2
8. time -> 0.98s

## Random Forest :

1. accuracy -> 85.47 | This means that the decision tree has correctly predicted the class on the test data 94.09% of the time
2. precision -> 90.60 | The model gave correct predictions for a class 1, the model predicted correctly
3. precision -> 62.77 | The model gave correct predictions for a class 2, the model predicted correctly
4. recall -> 0.9 | This is the fraction of instances of a class 1 that were correctly predicted, that is 0.9
5. recall -> 0.6 | This is the fraction of instances of a class 2 that were correctly predicted, that is 0.9
6. f1 -> 91.05 | This is the harmonic mean of precision and recall, for class 1
7. f1 -> 61.39 | This is the harmonic mean of precision and recall, for class 2
8. time -> 88.98s

Since we can see that we did not get any dramatic change while using random forest and the difference in time is huge. We can use random forest when we are suffering with overfitting. In this case, random forest did slightly better than decision tree.